

ECE 408 Final Project Report

Team Name: tbd

Yiqing Zhou (yiqing2)

Nuocheng Pan(np9)

April 17th, 2019

Milestone 1: RAI Setup

Kernels that collectively consume more than 90% of the program time:

```

40.45%:    [CUDA memcpy HtoD]
20.32%:    implicit_convolve_sgemm
11.88%:    volta_cgemm_64x32_tn
7.07%:     op_generic_tensor_kernel
5.62%:     volta_sgemm_128x128_tn
5.61%:     fft2d_c2r_32x32
4.52%:     pooling_fw_4d_kernel
3.70% :    fft2d_r2c_32x32

```

CUDA API calls that collectively consume more than 90% of the program time:

```

42.61%    cudaStreamCreateWithFlags
34.35%    cudaMemGetInfo
21.02%    cudaFree

```

Explanation of difference between kernels and API calls:

Kernels are functions programmed by users. Kernels are launched by host and run on devices. APIs are provided by CUDA runtime system and could be directly called by users.

CPU output and runtime: (runtime is bolded)

```

Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8236}
8.98user 3.57system 0:05.07elapsed 247%CPU (0avgtext+0avgdata
2470144maxresid
ent)k
0inputs+2824outputs (0major+668695minor)pagefaults 0swaps

```

GPU output and runtime: (runtime is bolded)

```

Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8236}
4.40user 3.12system 0:04.38elapsed 171%CPU (0avgtext+0avgdata
2840696maxresident)k
0inputs+4552outputs (0major+660254minor)pagefaults 0swaps

```

Milestone 2: CPU Convolution Implementation

OP and Exec Time for different input data size:

*** Running /usr/bin/time python m2.1.py 100**

```
Loading fashion-mnist data... done
Loading model... done
New Inference
Op Time: 0.034094
Op Time: 0.075474
Correctness: 0.84 Model: ece408
2.87user 2.76system 0:01.00elapsed 562%CPU (0avgtext+0avgdata
203620maxresident)k
0inputs+8outputs (0major+61034minor)pagefaults 0swaps
```

*** Running /usr/bin/time python m2.1.py 1000**

```
Loading fashion-mnist data... done
Loading model... done
New Inference
Op Time: 0.245769
Op Time: 0.749210
Correctness: 0.852 Model: ece408
4.29user 3.00system 0:02.00elapsed 363%CPU (0avgtext+0avgdata
331980maxresident)k
0inputs+2824outputs (0major+110686minor)pagefaults 0swaps
```

*** Running /usr/bin/time python m2.1.py 10000**

```
Loading fashion-mnist data... done
Loading model... done
New Inference
Op Time: 2.446601
Op Time: 7.594124
Correctness: 0.8397 Model: ece408
15.54user 4.46system 0:11.65elapsed 171%CPU (0avgtext+0avgdata
1617164maxresident)k
0input
s+2824outputs (0major+617305minor)pagefaults 0swaps
```

Milestone 3: GPU Forward Convolution

Nvprof result:

```
==278== NVPROF is profiling process 278, command: python m3.1.py
Loading model... done
New Inference
```

Op Time: 0.005826

Op Time: 0.031620

Correctness: 0.8397 Model: ece408

==278== Profiling application: python m3.1.py

==278== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	60.09%	37.378ms	2	18.689ms	5.7804ms	31.598ms	
mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)							
	27.21%	16.927ms	20	846.34us	1.1200us	16.409ms	[CUDA memcpy HtoD]
	3.96%	2.4646ms	2	1.2323ms	20.864us	2.4438ms	
volta_sgemv_32x128_tn							
	3.93%	2.4415ms	2	1.2207ms	737.50us	1.7040ms	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>,							
mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul,							
mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float,							
int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)							
	2.62%	1.6307ms	2	815.34us	22.368us	1.6083ms	void
op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7,							
cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*,							
cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float,							
float, dimArray, reducedDivisorArray)							
	1.70%	1.0576ms	1	1.0576ms	1.0576ms	1.0576ms	void
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *,							
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float,							
cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float,							
cudnnPoolingStruct, int, cudnn::reduced_divisor, float)							
	0.25%	152.89us	1	152.89us	152.89us	152.89us	void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,							
mshadow::Shape<int=2>, int=2, int)							
	0.12%	72.416us	1	72.416us	72.416us	72.416us	void
mshadow::cuda::SoftmaxKernel<int=8, float,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu,							
int=2, unsigned int)							
	0.04%	27.614us	13	2.1240us	1.1520us	6.5280us	void
mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int,							
mshadow::Shape<int=2>, int=2)							
	0.04%	23.711us	2	11.855us	2.4000us	21.311us	void
mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1,							
float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>,							
int=2)							
	0.02%	11.712us	10	1.1710us	992ns	1.5680us	[CUDA memset]
	0.01%	7.6160us	1	7.6160us	7.6160us	7.6160us	[CUDA memcpy DtoH]
	0.01%	4.9920us	1	4.9920us	4.9920us	4.9920us	void
mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,							
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,							
mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum,							

```

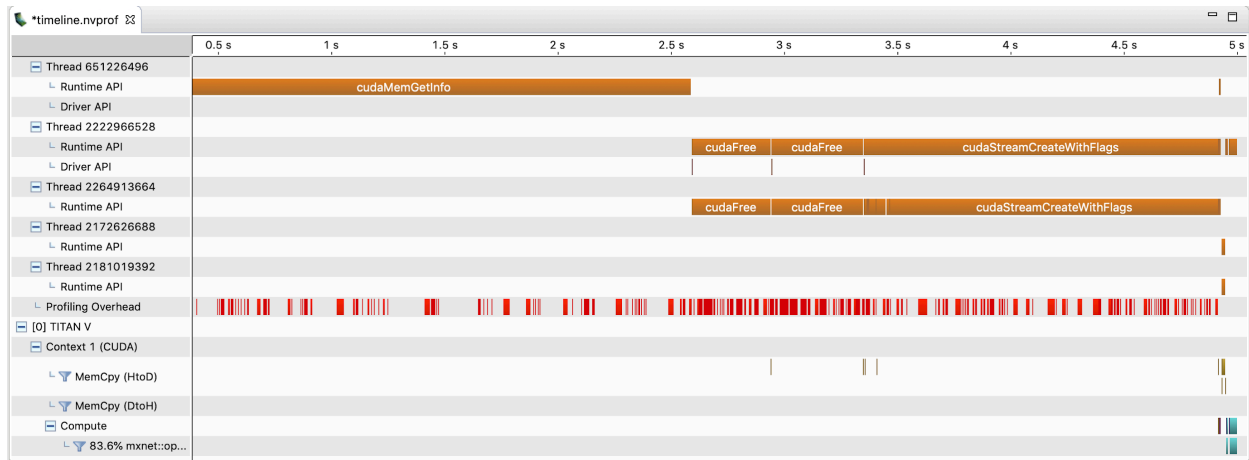
mshadow::Tensor<mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
    API calls: 41.86% 2.98971s 22 135.90ms 13.172us 1.56053s
cudaStreamCreateWithFlags
    33.80% 2.41345s 22 109.70ms 95.187us 2.40157s cudaMemGetInfo
    21.03% 1.50207s 18 83.449ms
841ns 399.12ms cudaFree
    1.13% 80.808ms 912 88.604us 305ns 23.436ms
cudaFuncSetAttribute
    0.62% 43.958ms 216 203.51us 855ns 25.633ms
cudaEventCreateWithFlags
    0.56% 39.842ms 6 6.6404ms 2.2730us 31.600ms
cudaDeviceSynchronize
    0.48% 34.204ms 9 3.8005ms 27.692us 16.468ms
cudaMemcpy2DAsync
    0.14% 9.9121ms 66 150.18us 5.5080us 2.0783ms cudaMalloc
    0.11% 8.1087ms 4 2.0272ms 499.17us 2.6772ms
cudaGetDeviceProperties
    0.09% 6.1239ms 12 510.33us 6.6460us 5.1190ms cudaMemcpy
    0.07% 4.7508ms 29 163.82us 2.1090us 2.2024ms
cudaStreamSynchronize
    0.04% 2.6436ms 375 7.0490us 272ns 334.18us
cuDeviceGetAttribute
    0.01% 909.36us 4 227.34us 45.866us 691.55us
cuDeviceGetName
    0.01% 887.76us 8 110.97us 13.114us 688.13us
cudaStreamCreateWithPriority
    0.01% 726.15us 2 363.08us 52.207us 673.94us cudaHostAlloc
    0.01% 713.23us 10 71.323us 7.5440us 480.63us
cudaMemsetAsync
    0.01% 646.22us 4 161.55us 92.941us 275.56us
cuDeviceTotalMem
    0.01% 617.82us 4 154.46us 74.072us 259.14us
cudaStreamCreate
    0.01% 540.37us 27 20.013us 8.2670us 52.173us
cudaLaunchKernel
    0.00% 292.81us 202 1.4490us 545ns 4.5470us
cudaDeviceGetAttribute
    0.00% 150.39us 29 5.1850us 1.0190us 16.214us cudaSetDevice
    0.00% 133.47us 6 22.244us 1.1630us 86.376us
cudaEventCreate
    0.00% 113.29us 557 203ns 75ns 812ns
cudaGetLastError
    0.00% 44.747us 18 2.4850us 581ns 4.4910us cudaGetDevice
    0.00% 26.923us 2 13.461us 5.0500us 21.873us
cudaHostGetDevicePointer
    0.00% 16.492us 4 4.1230us 1.7760us 7.1760us
cudaEventRecord
    0.00% 7.4790us 2 3.7390us 2.7320us 4.7470us cudaEventQuery
    0.00% 6.1870us 20 309ns 140ns 599ns
cudaPeekAtLastError
    0.00% 5.9340us 2 2.9670us 1.7720us 4.1620us
cudaDeviceGetStreamPriorityRange
    0.00% 5.4820us 6 913ns 441ns 1.8620us
cuDeviceGetCount
    0.00% 5.2120us 5 1.0420us 505ns 1.5460us cuDeviceGet
    0.00% 4.4920us 3 1.4970us 793ns 2.6370us cuInit

```

	0.00%	3.6930us	1	3.6930us	3.6930us	3.6930us
cuDeviceGetPCIBusId						
	0.00%	2.7100us	4	677ns	328ns	1.3210us
cuDeviceGetUuid						
	0.00%	2.0830us	3	694ns	338ns	1.3080us
cuDriverGetVersion						
	0.00%	1.9460us	4	486ns	239ns	806ns
cudaGetDeviceCount						

Milestone 4: GPU Forward Convolution Optimizations and Analysis

NVVP Timeline:



Optimization 1: Weight matrix (kernel values) in constant memory

Description: We observed that elements in kernel `k` are accessed by different threads in grid multiple times. By loading the kernel into the constant memory, we could decrease the number of global reads.

Code Snippet:

```
__constant__ float k_const[4096];  
...  
cudaMemcpyToSymbol(k_const,w.dptr_, sizeof(float)*C*K*K*M, 0);  
// access k from k_const in kernel function
```

Analysis:

Before this optimization, our kernel performance was dominantly limited by memory bandwidth (as shown in Figure 1.1). Before optimization, the op time data was (0.006040,0.031286). After optimization, we reduce it to (0.010241, 0.024736).

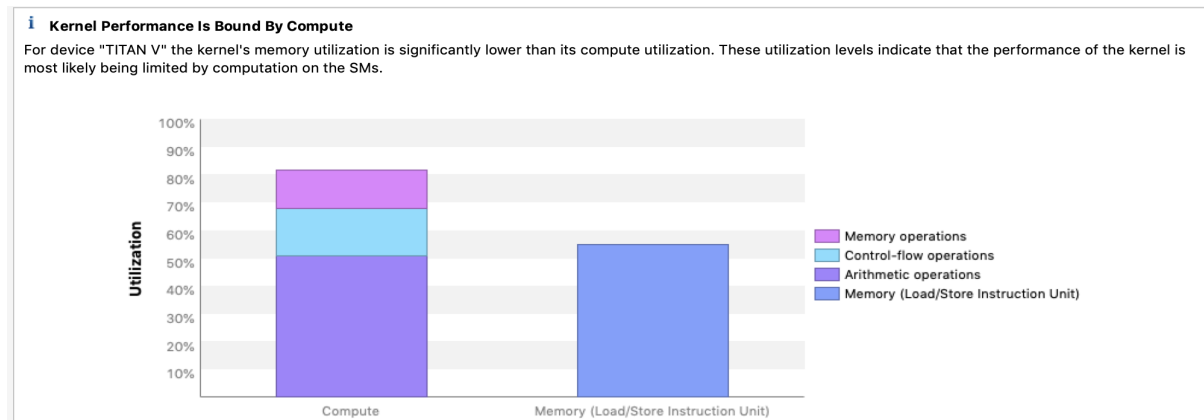


Figure 1.1 No optimization

Optimization 2: Shared Memory convolution

Description: In this optimization, we loaded the input x into shared memory. Asymptotically, this optimization should decrease global memory read by a factor of `TILE_WIDTH`. We tried two version of shared memory convolution.

Version 1 Code Snippet:

```
for (int c = 0; c < C; ++c){
    int s1 = threadIdx.y;
    while( s1 < TWK){
        int s2 = threadIdx.x;
        while( s2 < TWK){
            x_shared3d(c,s1,s2) = x4d(b,c, h + s1 - ty, w + s2 -
tx);

            s2 += TILE_WIDTH;
        }
        s1 += TILE_WIDTH;
    }
}
```

Analysis:

We use threads to put the input elements needed in to the shared memory of current block. We observed that in the same block, threads with same `Idx x` tends to repeatedly load some of elements that overlaps and is consecutive. Using shared memory may save some time by reducing global read and may utilize some of coalescing.

Version 2 Snippet:

```
for(int c=0; c<C; ++c){
    int tid = ty*TILE_WIDTH+tx;
    while(tid<TWK*TWK){
        // if(h+(tid/TWK)-ty<H && w+tid%TWK-tx<W)
        x_shared[c*TWK*TWK+tid] = x4d(b,c,h+(tid/TWK)-
ty,w+(tid%TWK)-tx);
        tid+=TILE_WIDTH*TILE_WIDTH;
    }
}
```

Analysis: Version 1 involves lots of control divergence due to the way we map threads onto the 2d array we want to load. Version 2 aims at reducing control divergence and taking advantage of memory coalescing.

From nvvp profile, we could see control divergence is drastically decreased after this modification. Following are some graph from NVVP.

Comparing 2 Versions:

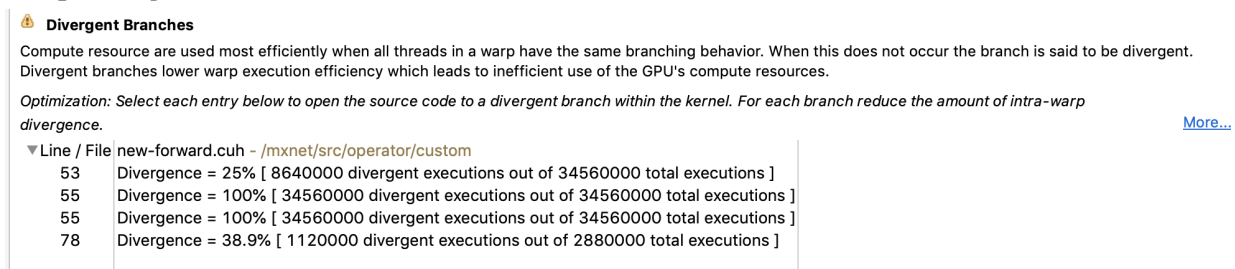


Figure 2.1: Optimization 1 and Version 1 of Optimization 2

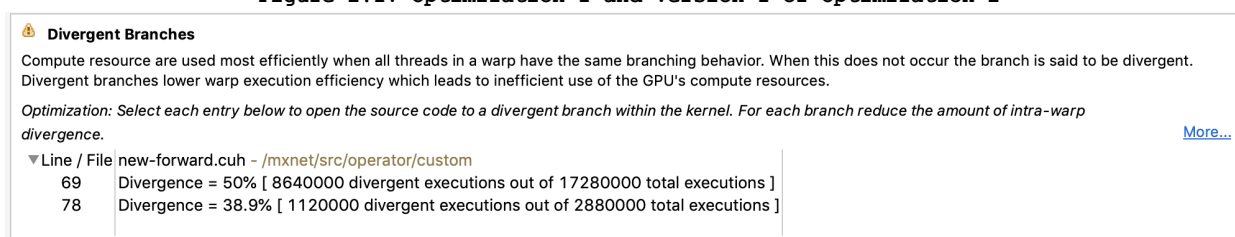


Figure 2.2: Optimization 1 and Version 2 of Optimization 2

Comment: Loading k and x from global memory to shared memory involves boundary checks, which will lead to control divergence. As shown in figure 2 and 3, adding version 1 of optimization 2 leads to a high control divergence rate. Comparing figure 3 and 4, Version 2 effectively decreases control divergence rate.

Optimization 3: Sweeping various parameters to find best values for block size

Description:

Based on the previous two optimizations, we tried different block sizes to get even better performance. With optimization 1 and 2

implemented, we noticed that our kernel performance was limited by computation on the SMs, which means when memory are read, SMs do not have enough computational resources. Thus we decrease number of threads in each block to resolve this problem. As shown in Figure 3.1, when using block dimension 16×16 , the performance of the kernel suffer from insufficient computational resources on SMs. For block dimension of 8×8 , it is no longer the case.

i Kernel Performance Is Bound By Compute

For device "TITAN V" the kernel's memory utilization is significantly lower than its compute utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by computation on the SMs.




Figure 3.1 Block dimension 16×16

i Kernel Performance Is Bound By Compute And Memory Bandwidth

For device "TITAN V" compute and memory utilization are balanced. These utilization levels indicate that kernel performance is good, but that additional performance improvement may be possible if either of both of compute and memory utilization levels are increased.




Figure 3.2 Block dimension 8×8

We could also see the improvement in performance simply from the decrease in the bolded operation time.

Output of Size of 32×32 in 10000 data size:

New Inference

Op Time: 0.014388

Op Time: **0.034988**

Correctness: 0.8397 Model: ece408

4.40user 3.34system 0:04.44elapsed 174%CPU (0avgtext+0avg
data 2856860maxresident)k
0inputs+4640outputs (0major+666082minor)pagefaults 0swaps

Output of Size of 16*16 in 10000 data size:

New Inference

Op Time: 0.006355

Op Time: **0.031710**

Correctness: 0.8397 Model: ece408

4.43user 3.32system 0:04.43elapsed 175%CPU (0avgtext+0avgdata
2844352maxresident)k

0inputs+4640outputs (0major+662779minor)pagefaults 0swaps

Output of Size of 8*8 in 10000 data size:

New Inference

Op Time: 0.010416

Op Time: **0.026588**

Correctness: 0.8397 Model: ece408

4.54user 2.97system 0:04.59elapsed 163%CPU (0avgtext+0avgdata
2858084maxresident)k

0inputs+4640outputs (0major+666778minor)pagefault
s 0swaps

Analysis:

Potential

Running in different block size, we found for current optimization, size of 8*8 is optimal. In most obvious way, it uses least time in completing second layer, so it is the best, but this may change subject to different opt implemented in final run. From figure 3.1 and 3.2 we can see the reason: Both has high computational power utilization but 16*16 has a lower memory usage, the performance is most limited by compute. Thus in 8*8, by decreasing the size of block, we increase computation and memory usage ratio so none of them remain idle while another is doing work.